

Autonomous Drone Racing Optimization vs. RL

Autonomous Drone Racing - Final Presentation
17.07.2024

Tim Lindenau




The Challenge – Autonomous Drone Racing In Uncertain Environments



Goal: Developing high level controller to safely slalom through a set of gates in the shortest time possible

Challenge: Uncertainties in the position of gates and obstacles require within-flight updates as true gate position only becomes available shortly before passing



Two Possible Approaches: Optimization and Reinforcement Learning



Approach 1: Optimization

Approach 2: Reinforcement Learning

Optimization Based Racing- General Approach

Block Diagram – Optimization Based Drone Racing

```
graph LR; Input[Nominal Gate and Obstacle Position] --> PP[Path Planning]; PP --> TG[Trajectory Generation]; subgraph TG; PS[Path Smoothing]; TP[Time Parametrization]; end; TG --> RP[Re-Planning]; RP --> PP; RP -->|Gate update| Input; RP -->|Drone position| Input;
```

Step 1: Path Planning

- **Goal:** Identify time-minimum path passing all gates
- **Minimum time** path depends on **dynamic limits** and **infeasible** to find with **real-time** constraints
- **Minimum path length** utilized as **proxy** due to **high computational efficiency** (<100ms)

Step 2: Trajectory Generation

- Find **trajectory**, i.e. mapping from time to control input
 $f: [0, T] \rightarrow [x, y, z, \dot{x}, \dot{y}, \dot{z}, \ddot{x}, \ddot{y}, \ddot{z}]$
- **Iterative path smoothing** minimizes discontinuities, while guaranteeing **path validity** after smoothing [1]
- **Forward and backward integration** efficient algorithm to find optimal time-parametrization s.t. drones' velocity and acceleration constraints [2]

7

Action and Observation Space Enable Flight With Disturbances

Two factors for designing effective observation and action space

- 1 Simplify spaces** as much as possible
→ **Do not confuse** learning
- 2 Provide gate and obstacle information in local frame**
→ **Globally generalizing** behavior

Observation and action space transformation pipeline

Position	Velocity	Orientation	Body Rates	Gates	Obstacles
x, y, z	v_x, v_y, v_z	ϕ, θ, ψ	$\omega_\phi, \omega_\theta, \omega_\psi$	G_0, G_1, \dots	O_0, O_1, \dots
Simplify Observation Space 1					
x, y, z	v_x, v_y, v_z	ϕ, θ, ψ	$\omega_\phi, \omega_\theta, \omega_\psi$	G_0, G_1, \dots	O_0, O_1, \dots
Translate to local frame 2					
x, y, z	v_x, v_y, v_z	ϕ, θ	$\omega_\phi, \omega_\theta$	G_1	O_1
Policy Network					
$\Delta x, \Delta y, \Delta z$					
De-normalize Action 2					
x, y, z					
Control Command					

Legend: ● Global frame ● Local frame

11

Tim Lindenau

Two Possible Approaches: Optimization and Reinforcement Learning



Approach 1: Optimization

Approach 2: Reinforcement Learning

Optimization Based Racing- General Approach

Block Diagram – Optimization Based Drone Racing

```
graph LR; Input[Nominal Gate and Obstacle Position] --> PathPlanning[Path Planning]; PathPlanning --> TrajectoryGeneration[Trajectory Generation]; subgraph TrajectoryGeneration; PathSmoothing[Path Smoothing]; TimeParametrization[Time Parametrization]; end; TrajectoryGeneration --> RePlanning[Re-Planning]; RePlanning --> PathPlanning; RePlanning --> RePlanning; GateUpdate[Gate update] --> RePlanning; DronePosition[Drone position] --> RePlanning;
```

Step 1: Path Planning

- **Goal:** Identify time-minimum path passing all gates
- **Minimum time** path depends on **dynamic limits** and **infeasible** to find with **real-time** constraints
- **Minimum path length** utilized as **proxy** due to **high computational efficiency** (<100ms)

Step 2: Trajectory Generation

- Find **trajectory**, i.e. mapping from time to control input
 $f: [0, T] \rightarrow [x, y, z, \dot{x}, \dot{y}, \dot{z}, \ddot{x}, \ddot{y}, \ddot{z}]$
- **Iterative path smoothing** minimizes discontinuities, while guaranteeing **path validity** after smoothing [1]
- **Forward and backward integration** efficient algorithm to find optimal time-parametrization s.t. drones' velocity and acceleration constraints [2]

7

Action and Observation Space Enable Flight With Disturbances

Two factors for designing effective observation and action space

- 1 **Simplify spaces** as much as possible
→ **Do not confuse** learning
- 2 **Provide gate and obstacle information in local frame**
→ **Globally generalizing** behavior

Observation and action space transformation pipeline

Position	Velocity	Orientation	Body Rates	Gates	Obstacles
x, y, z	v_x, v_y, v_z	ϕ, θ, ψ	$\omega_x, \omega_y, \omega_z$	G_0, G_1, \dots	O_0, O_1, \dots
Simplify Observation Space 1					
x, y, z	v_x, v_y, v_z	ϕ, θ, ψ	$\omega_x, \omega_y, \omega_z$	G_0, G_1, \dots	O_0, O_1, \dots
Translate to local frame 2					
x, y, z	v_x, v_y, v_z	ϕ, θ, ψ	$\omega_x, \omega_y, \omega_z$	G_0, G_1, \dots	O_0, O_1, \dots
Policy Network					
$\Delta x, \Delta y, \Delta z$					
De-normalize Action 3					
x, y, z					
Control Command					

11

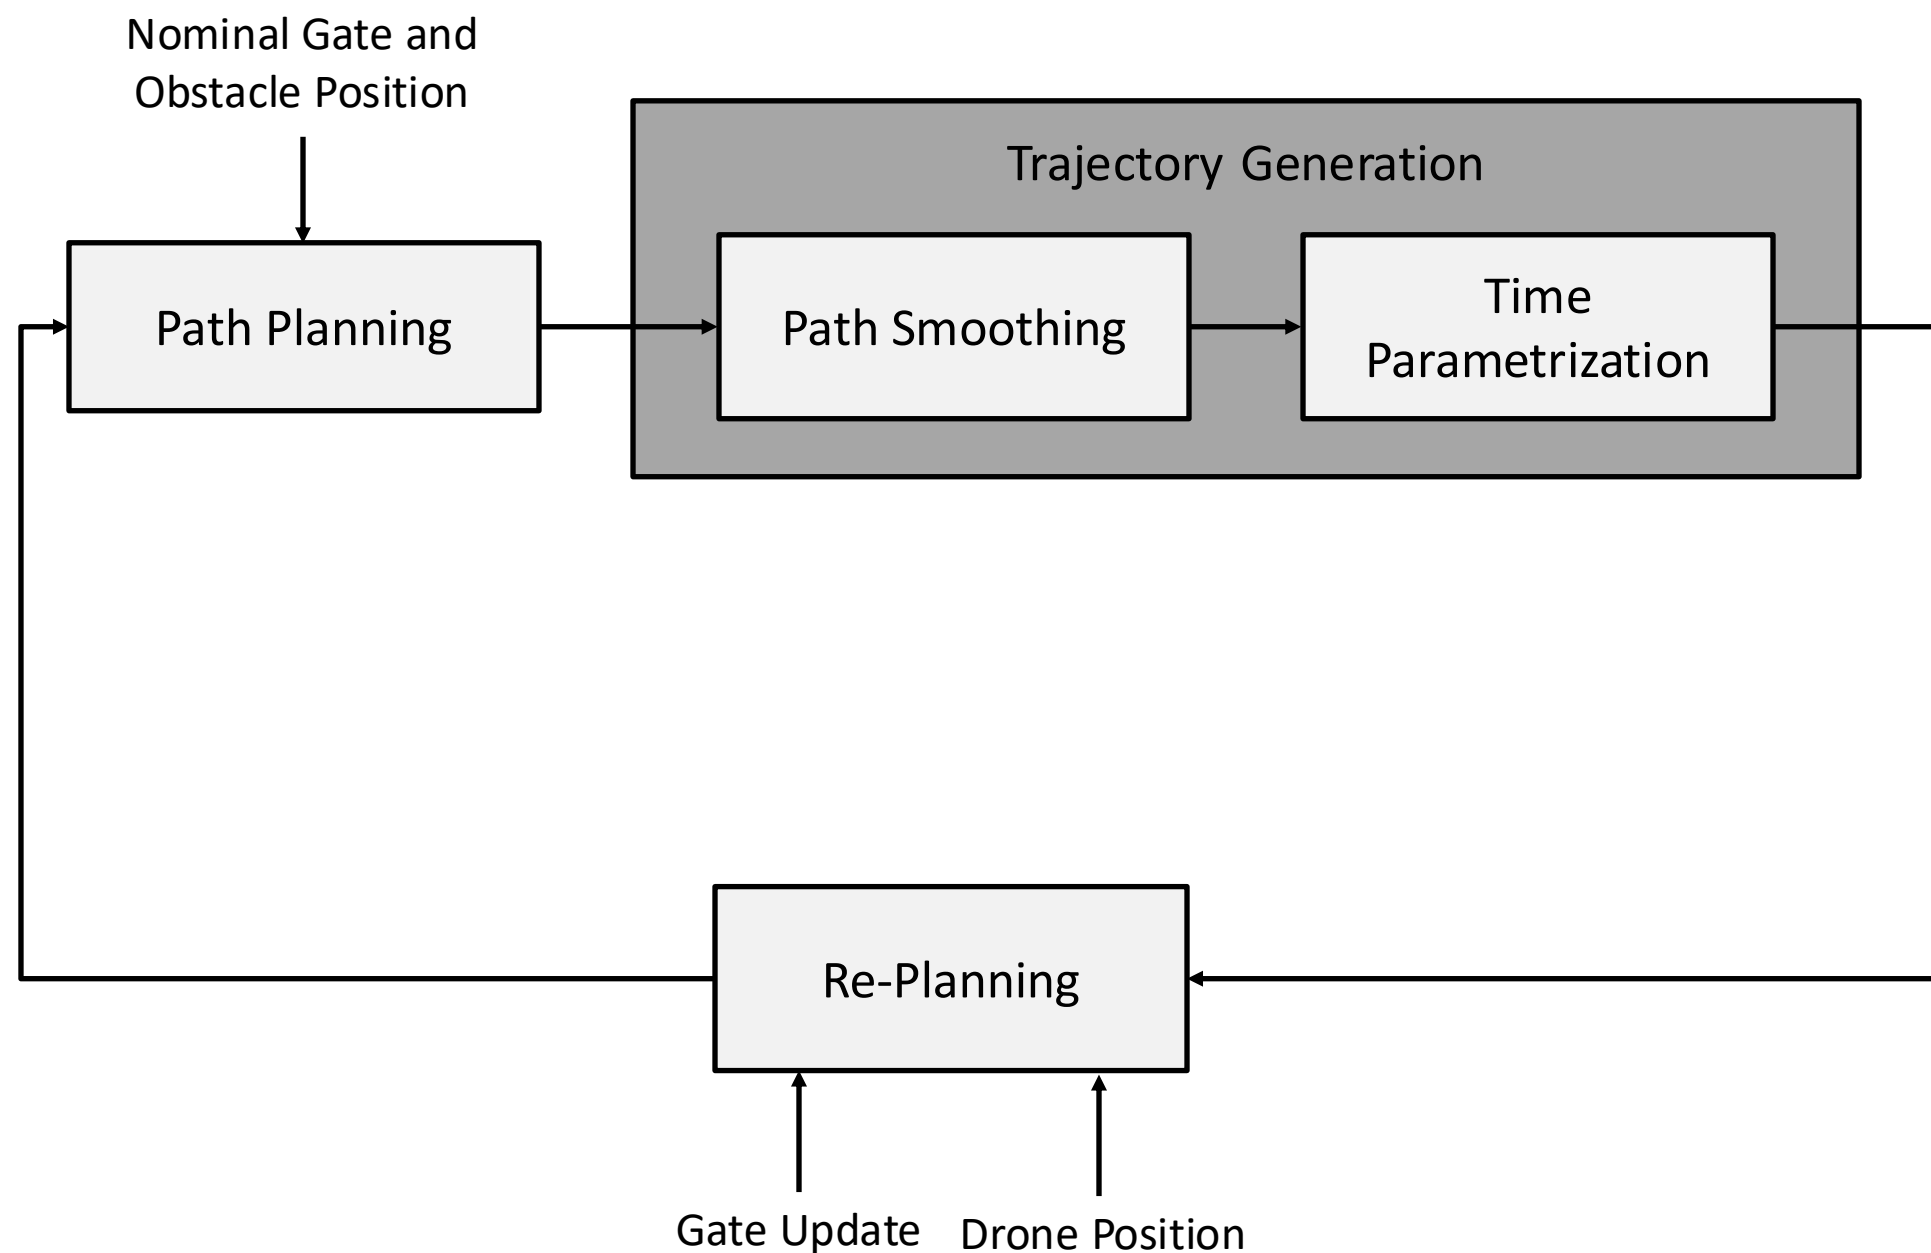
Tim Lindenau

Global frame Local frame

Optimization Based Racing- General Approach



Block diagram – Optimization based drone racing



Step 1: Path planning

- **Goal:** Identify time-minimum path
- **Minimum time** path depends on **dynamics** and **computationally expensive** to find
- **Minimum path length** utilized as **proxy** due to **high computational efficiency** (<100ms)

Step 2: Trajectory generation

- Find **trajectory**
 $f: [0, T] \rightarrow [x, y, z, \dot{x}, \dot{y}, \dot{z}, \ddot{x}, \ddot{y}, \ddot{z}]$
- **Iterative path smoothing** minimizes discontinuities, while guaranteeing **path validity** [1]
- **Forward and backward integration**¹ to find **optimal time-parametrization** s.t. drones' velocity and acceleration constraints [2]

Step 3: Re-planning

- Re-plan **within flight** to adapt to updates
- Currently **only updating gates** not obstacles



Re-planning: The challenge

- 1 **Computational efficient** trajectory generation pipeline, enables computation times of around **100ms**
- 2 **Too slow** to enable within-flight re-computations. **Drone moving** multiple centimeters **without control input**

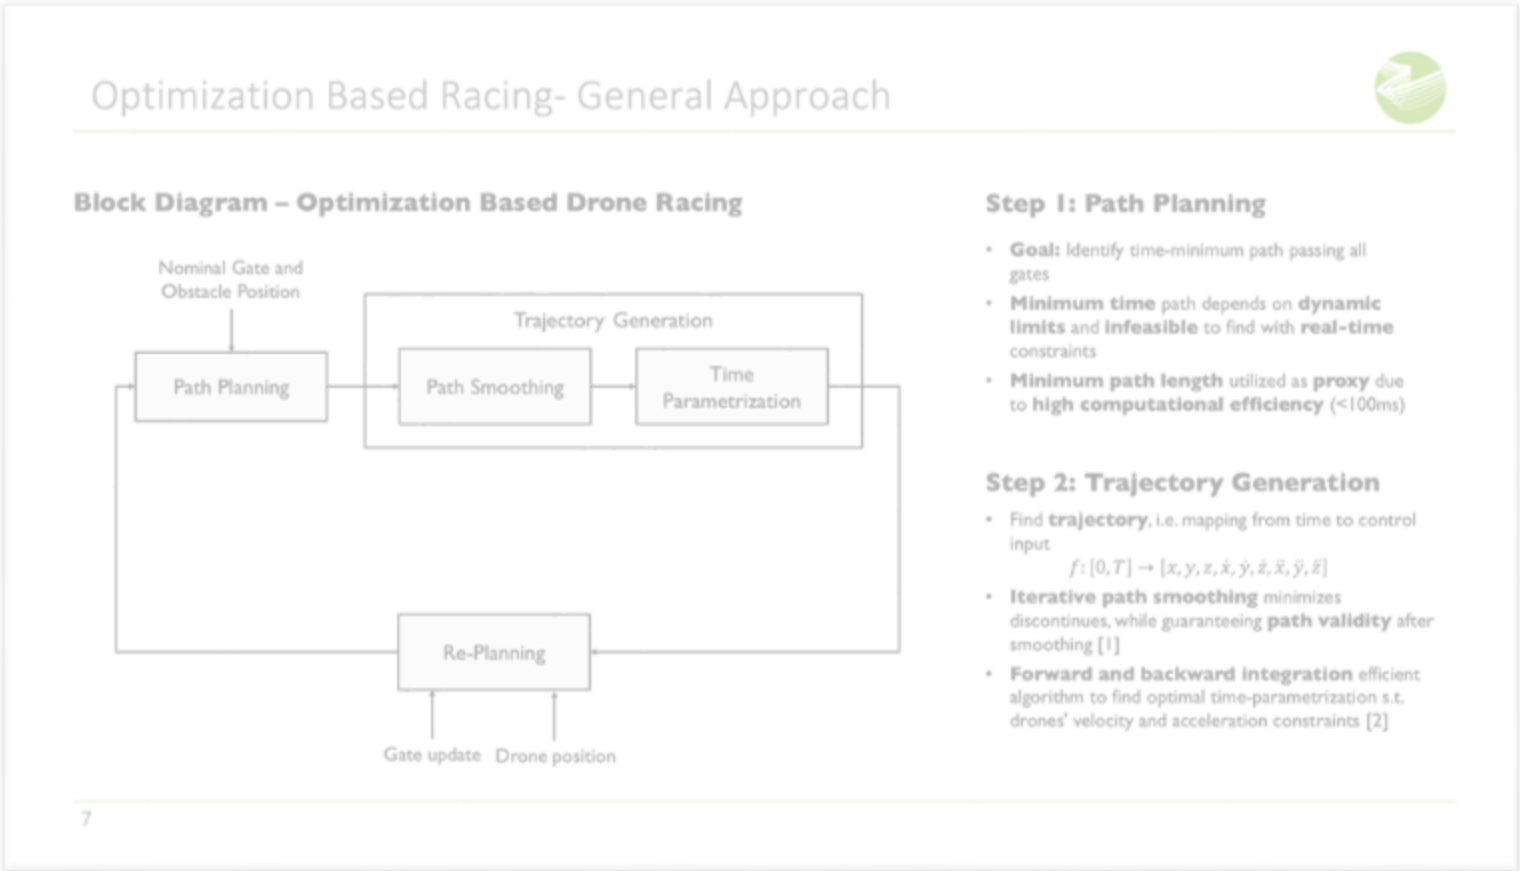
Non-blocking implementation and start-point offset enable re-planning without loss of control

- 1 **Non-blocking trajectory generation**
 - Trajectory computation in **extra thread** prevents **blocking** of **main control loop**
 - **Controller** continues based on **previous trajectory** until re-planning completed
- 2 **Start-point offset**
 - **Start point offsetted** along trajectory, based on expected computation time
 - Trajectories **smoothly merged** after re-planning

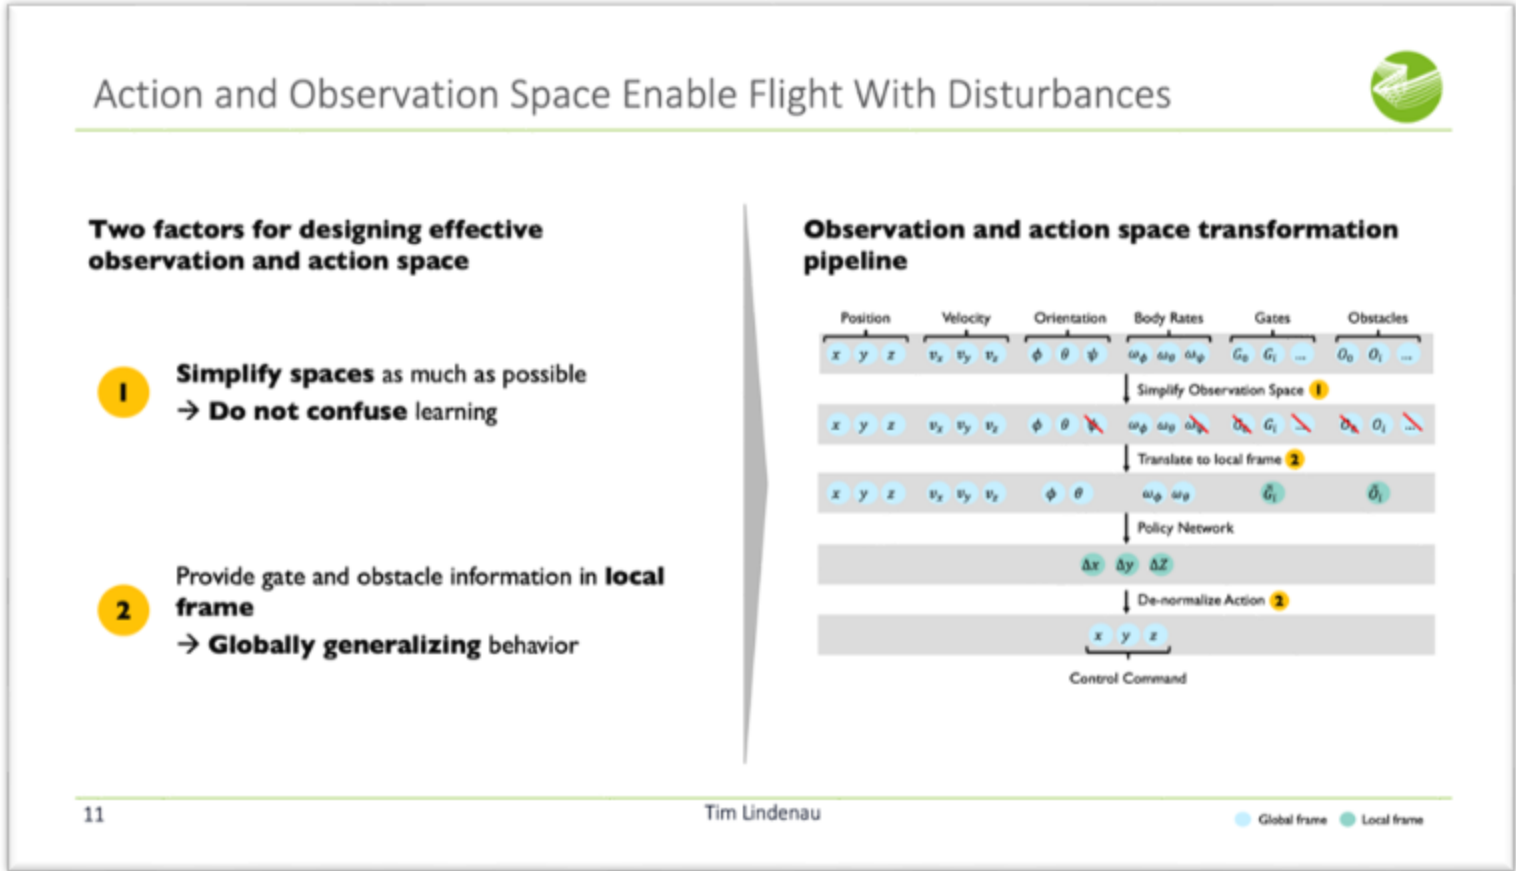
Two Possible Approaches: Optimization and Reinforcement Learning



Approach 1: Optimization



Approach 2: Reinforcement Learning





Reinforcement learning: Motivation & challenge

Generalization to disturbances enabled by well designed action, -and observation space

1

Optimization complex and with **tradeoffs** in flight time to enable re-planning capabilities

2

Reinforcement learning outperformed optimization in previous work [3, 4]

3

Weak generalization capabilities, potentially **hinder adapting** to **changing gate positions**

Action and Observation Space Enable Flight With Disturbances

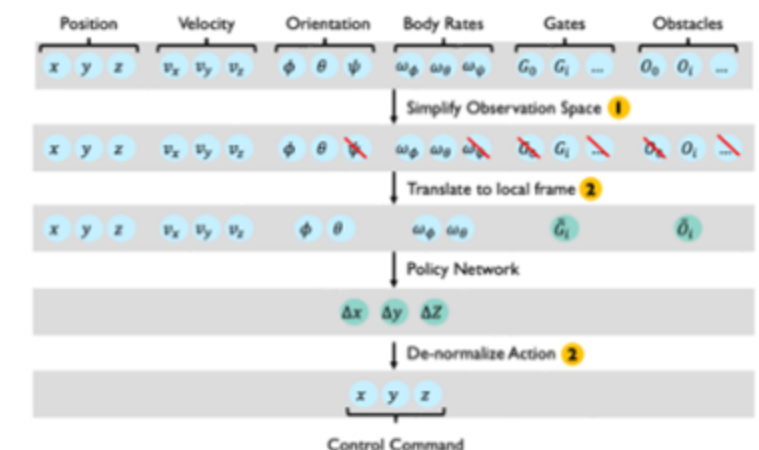


Two factors for designing effective observation and action space

1 **Simplify spaces** as much as possible
→ **Do not confuse** learning

2 Provide gate and obstacle information in **local frame**
→ **Globally generalizing** behavior

Observation and action space transformation pipeline



11

Tim Lindenau

Reward function also important.
However good results with common
per time-step progress



Two factors for designing effective observation and action space

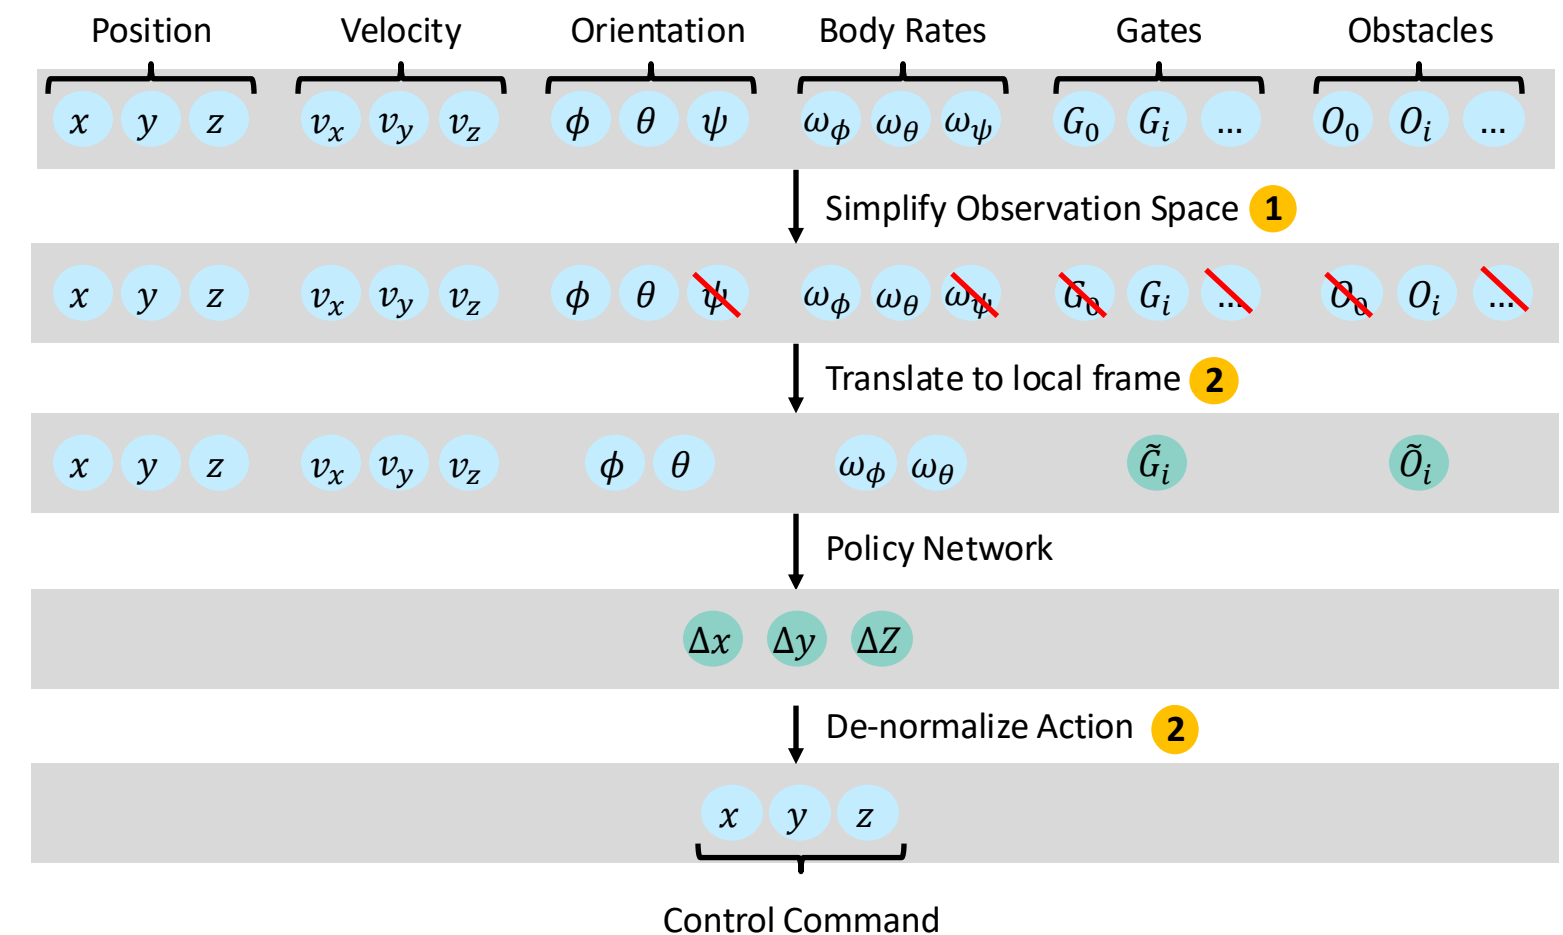
1

Simplify spaces as much as possible
→ **Do not confuse** learning

2

Provide gate and obstacle information in **local** frame
→ **Globally generalizing** behavior

Observation and action space transformation pipeline



Experimental Results (Simulation)



Result collection – Average lap times for different approaches
[values collected over 100 runs]

Algorithm	Difficulty	Mean Time (s)	Success Rate (%)
Baseline	Static	12.47	100
Baseline	Disturbed	12.47	18
Optimization	Static	6.50	86
RL	Static	5.75	100
Optimization	Disturbed	9.78	81
RL	Disturbed	6.49	47

1

Optimization and RL, twice as fast as baseline in static environment without disturbances

2

Baseline without re-planning **fails** in **disturbed environment**, only reaching success rate of 18%

3

Optimization most robust in disturbed environment, however **tradeoff in lap time** required to reach robustness


4

50% success rate in disturbed environments indicates **RL agent learning to react to shifted gates**

Three Factors Enable RL to Outperform Optimization in Lap-Time



Flight trajectories: Optimization (blue) vs. Reinforcement Learning (orange)



Three factors explain how the RL agent can outperform optimization in lap-time

- 1** **Missing controllability** due to onboard controller indirection cause **overshoots**
- 2** **Shortest path length** criterion introduces sharp turns **reducing speed** and **amplifying overshoots**
- 3** **Padding around obstacles** required to guarantee robustness **increases path length**



1

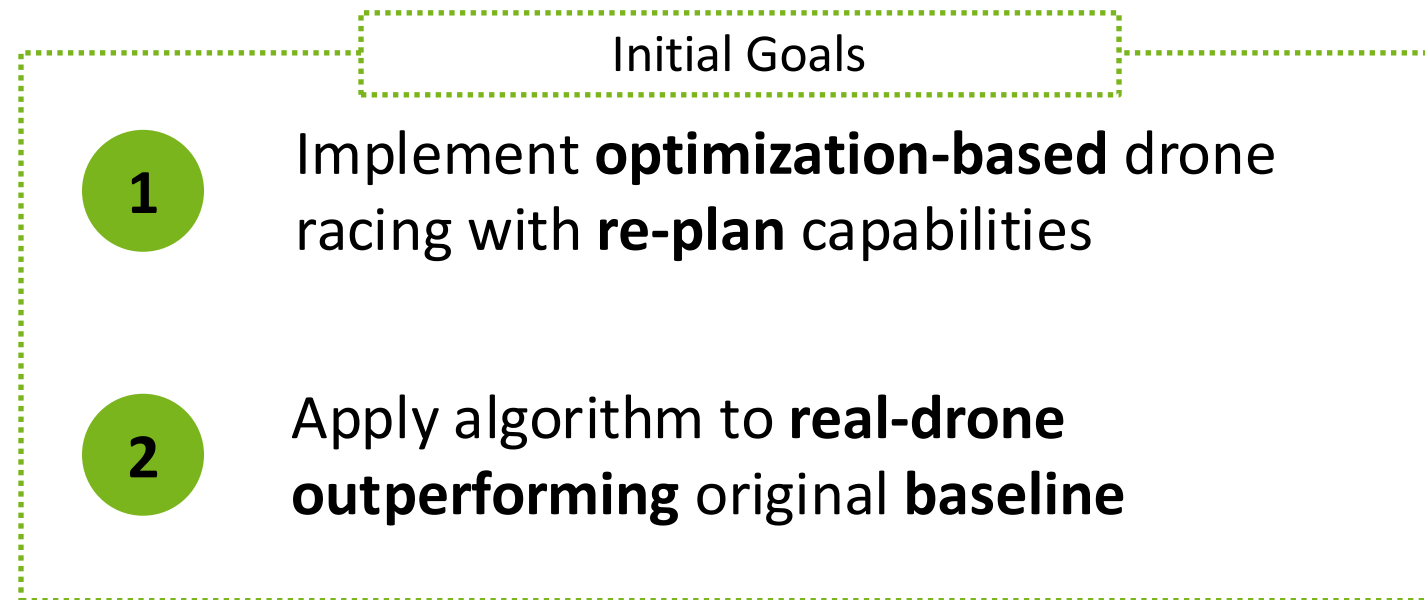
Optimization approach **successfully** transferred to **real world** (video)

2

No time to transfer RL approach. Experiences of other groups indicate **challenges**



Milestones



- 3 Develop **fast RL agent** for **static environment**
- 4 Develop **fast RL agent** for **disturbed environment**
- 5 **Sim2Real** transfer **RL agent**

Key Results

- **Achieved** optimization-based drone racing with re-planning by focusing on **computational efficient multithreaded C++** implementation
 - Reached **fast lap times in static** environment and **high robustness** to disturbances
 - Future work: Include **obstacle updates**
-
- **Solved static environment** with 100% success rates and fast lap times
 - **Tuning observation and action space** enabled flying in **dynamic environment**
 - Future work: Further **improve robustness** and **Sim2Real** transfer



1. “Ompl::Geometric::RRTstar Class Reference.” Accessed July 6, 2024.
https://ompl.kavrakilab.org/classompl_1_1geometric_1_1RRTstar.html.
2. Kunz, Tobias, and Mike Stilman. “Time-Optimal Trajectory Generation for Path Following with Bounded Acceleration and Velocity.” In *Robotics*, edited by Nicholas Roy, Paul Newman, and Siddhartha Srinivasa, 209–16. The MIT Press, 2013. <https://doi.org/10.7551/mitpress/9816.003.0032>.
3. “Champion-Level Drone Racing Using Deep Reinforcement Learning | Nature.” Accessed April 23, 2024.
<https://www.nature.com/articles/s41586-023-06419-4>.
4. Song, Yunlong, Mats Steinweg, Elia Kaufmann, and Davide Scaramuzza. “Autonomous Drone Racing with Deep Reinforcement Learning.” In *2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 1205–12. Prague, Czech Republic: IEEE, 2021.
<https://doi.org/10.1109/IROS51168.2021.9636053>.

Appendix


Autonomous Drone Racing - Final Presentation
17.07.2024

Tim Lindenau






Approach 1: Optimization



Approach 2: Reinforcement Learning





Key Requirements For Effective Reward Function

Path progress provides dense proxy reward for minimizing lap times

- 1 **Alignment** with goal of **drone racing**, i.e., finishing track in minimum time
- 2 **Dense supervision** at each time-step to simplify information gain of learning algorithm

- **Lap time reward** immediately rewards fast lap time, however reward sparse only at end of successful episodes
- **Per-time-step progress** towards next goal, good **proxy** reward, **providing dense supervision** at each time step

$$r_{prog}(t) = \|\mathbf{x}(t-1) - \mathbf{g}\|_2 - \|\mathbf{x}(t) - \mathbf{g}\|_2$$

- **Total reward function** completed with binary gate-pass, and collision reward

$$r(t) = \lambda_{prog} r_{prog}(t) + \lambda_{gate} r_{gate}(t) + \lambda_{col} r_{col}(t)$$